


dsc-phase1-project-final

April 1, 2021

1 Getting Started in the Movie Industry

Author: Jonathan Lee ***



Dvd Vectors by Vecteezy

1.1 Overview

This analysis focuses on an exploration of data tables from IMDB and The Numbers. We will walk through high-level exploratory data analysis to preview what data we have to work with, data cleansing to ensure that our data is accurate, feature engineering to calculate useful metrics based on the provided data and visualization to effectively explain what aspects of film-making Microsoft should focus on as it embarks on a journey to find success in the film industry. The results from this analysis show that overall, Musicals tend to perform well, but of course there are multiple other factors to keep in consideration.

1.2 Business Problem

Microsoft sees all the big companies creating original video content and they want to get in on the fun. They have decided to create a new movie studio, but they don't know anything about creating movies. The goal of this analysis is to explore what types of films are currently doing the best at the box office and translate those findings into actionable insights that the head of Microsoft's new movie studio can use to help decide what type of films to create. For this analysis, we will focus on the following questions:

1. What genres of movie are likely to succeed?
2. What genres cost more to produce?
3. How does production budget affect the success of a movie?

1.3 Data Understanding

In this analysis, we will be exploring datasets from IMDB and The Numbers. We can obtain information regarding movie titles, release years, genres, budget, and ratings from IMDB, and financial data from The Numbers.

To create an optimal basis for analysis, we will restrict movie releases from 2015 to 2019. Because movie preferences change with social context over generations, this will ensure that we have a dataset of movies that are relevant in this time period. By setting an upper limit for movie releases at 2019, we also eliminate any unusual data that has been impacted by COVID-19 restrictions. It is safe to make this assumption, since we are seeing more states open up restrictions with the decline of COVID-19 cases.

We will also be focusing on financial information within the US, since it would be best for a new film studio to focus on a specific audience as opposed to a worldwide audience.

1.3.1 Previewing the Data Tables

```
[1]: # Import standard packages.
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from matplotlib.ticker import FuncFormatter

%matplotlib inline
```

*include a note regarding how i chose my datasets

```
[2]: # Load and preview data files for ratings and revenue.

imdb_title_basics_df = pd.read_csv('zippedData/imdb.title.basics.csv.gz')
imdb_title_ratings_df = pd.read_csv('zippedData/imdb.title.ratings.csv.gz')
budgets_df = pd.read_csv('zippedData/tn.movie_budgets.csv.gz')

display(imdb_title_basics_df.head(),
        imdb_title_ratings_df.head(),
```

```
budgets_df.head()
```

	tconst	primary_title	original_title \
0	tt0063540	Sunghursh	Sunghursh
1	tt0066787	One Day Before the Rainy Season	Ashad Ka Ek Din
2	tt0069049	The Other Side of the Wind	The Other Side of the Wind
3	tt0069204	Sabse Bada Sukh	Sabse Bada Sukh
4	tt0100275	The Wandering Soap Opera	La Telenovela Errante

	start_year	runtime_minutes	genres
0	2013	175.0	Action, Crime, Drama
1	2019	114.0	Biography, Drama
2	2018	122.0	Drama
3	2018	NaN	Comedy, Drama
4	2017	80.0	Comedy, Drama, Fantasy

	tconst	averagerating	numvotes
0	tt10356526	8.3	31
1	tt10384606	8.9	559
2	tt1042974	6.4	20
3	tt1043726	4.2	50352
4	tt1060240	6.5	21

	id	release_date	movie \
0	1	Dec 18, 2009	Avatar
1	2	May 20, 2011	Pirates of the Caribbean: On Stranger Tides
2	3	Jun 7, 2019	Dark Phoenix
3	4	May 1, 2015	Avengers: Age of Ultron
4	5	Dec 15, 2017	Star Wars Ep. VIII: The Last Jedi

	production_budget	domestic_gross	worldwide_gross
0	\$425,000,000	\$760,507,625	\$2,776,345,279
1	\$410,600,000	\$241,063,875	\$1,045,663,875
2	\$350,000,000	\$42,762,350	\$149,762,350
3	\$330,600,000	\$459,005,868	\$1,403,013,963
4	\$317,000,000	\$620,181,382	\$1,316,721,747

1.3.2 Previewing the Data Types

```
[3]: # Display data types of each column in each table
```

```
display(imdb_title_basics_df.info(),
        imdb_title_ratings_df.info(),
        budgets_df.info())
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 146144 entries, 0 to 146143
Data columns (total 6 columns):
```

#	Column	Non-Null Count	Dtype
0	tconst	146144 non-null	object
1	primary_title	146144 non-null	object
2	original_title	146123 non-null	object
3	start_year	146144 non-null	int64
4	runtime_minutes	114405 non-null	float64
5	genres	140736 non-null	object

dtypes: float64(1), int64(1), object(4)

memory usage: 6.7+ MB

<class 'pandas.core.frame.DataFrame'>

RangeIndex: 73856 entries, 0 to 73855

Data columns (total 3 columns):

#	Column	Non-Null Count	Dtype
0	tconst	73856 non-null	object
1	averagerating	73856 non-null	float64
2	numvotes	73856 non-null	int64

dtypes: float64(1), int64(1), object(1)

memory usage: 1.7+ MB

<class 'pandas.core.frame.DataFrame'>

RangeIndex: 5782 entries, 0 to 5781

Data columns (total 6 columns):

#	Column	Non-Null Count	Dtype
0	id	5782 non-null	int64
1	release_date	5782 non-null	object
2	movie	5782 non-null	object
3	production_budget	5782 non-null	object
4	domestic_gross	5782 non-null	object
5	worldwide_gross	5782 non-null	object

dtypes: int64(1), object(5)

memory usage: 271.2+ KB

None

None

None

1.4 Data Preparation

Because the provided tables currently do not have much meaning by themselves, we need to address missing and duplicated data, and we also need to merge the tables in a way that preserves the accuracy of the data.

1.4.1 Merging IMDB Data

We begin by merging the two datasets pulled from IMDB on their common id key labeled “tconst” and addressing missing values for genres and average rating, which are both criteria that will be

importance in our final analysis.

```
[4]: # Merge imdb tables to pair movie titles with their ratings.
```

```
imdb_df = imdb_title_basics_df.merge(imdb_title_ratings_df, how='right',
                                     on='tconst')

print("number of rows:", len(imdb_df))
display(imdb_df.head())
```

number of rows: 73856

	tconst	primary_title	original_title	start_year	\
0	tt10356526	Laiye Je Yaarian	Laiye Je Yaarian	2019	
1	tt10384606	Borderless	Borderless	2019	
2	tt1042974	Just Inès	Just Inès	2010	
3	tt1043726	The Legend of Hercules	The Legend of Hercules	2014	
4	tt1060240	Até Onde?	Até Onde?	2011	

	runtime_minutes	genres	averagerating	numvotes
0	117.0	Romance	8.3	31
1	87.0	Documentary	8.9	559
2	90.0	Drama	6.4	20
3	99.0	Action,Adventure,Fantasy	4.2	50352
4	73.0	Mystery,Thriller	6.5	21

```
[5]: imdb_df.isna().sum()
```

```
[5]: tconst          0
primary_title      0
original_title     0
start_year         0
runtime_minutes    7620
genres             804
averagerating      0
numvotes           0
dtype: int64
```

```
[6]: # Fill missing genres with 'None' and filter out rows that are missing
# averagerating.
```

```
imdb_df['genres'].fillna('None', inplace=True)
imdb_df = imdb_df[imdb_df['averagerating'].notna()]
```

1.4.2 Converting Data Types

We want to ensure that our movie titles from the IMDB dataset match with the correct titles from the The Numbers dataset, so we will eventually merge them on the movie title and year. Before we

can do so, we need to convert the date information provided in The Numbers dataset into a type and format that matches the date information from IMDB.

```
[7]: # Convert release date to show year only as int type.

budgets_df['release_date'] = budgets_df['release_date'].str[-4:].astype(int)
```

We also need to format the dollar amounts in order to be able to correctly graph the financial data.

```
[8]: # Remove commas from dollar amounts.

budgets_df['production_budget'] = budgets_df['production_budget'] \
    .replace(',', '', regex = True)
budgets_df['domestic_gross'] = budgets_df['domestic_gross'] \
    .replace(',', '', regex = True)
budgets_df['worldwide_gross'] = budgets_df['worldwide_gross'] \
    .replace(',', '', regex = True)

# Remove $ sign and convert to int type.

budgets_df['production_budget'] = budgets_df['production_budget'] \
    .str[1:].astype(int)
budgets_df['domestic_gross'] = budgets_df['domestic_gross'] \
    .str[1:].astype(int)
budgets_df['worldwide_gross'] = budgets_df['worldwide_gross'].str[1:] \
    .astype(int)

budgets_df.head()
```

```
[8]:
```

	id	release_date	movie
0	1	2009	Avatar
1	2	2011	Pirates of the Caribbean: On Stranger Tides
2	3	2019	Dark Phoenix
3	4	2015	Avengers: Age of Ultron
4	5	2017	Star Wars Ep. VIII: The Last Jedi

	production_budget	domestic_gross	worldwide_gross
0	425000000	760507625	2776345279
1	410600000	241063875	1045663875
2	350000000	42762350	149762350
3	330600000	459005868	1403013963
4	317000000	620181382	1316721747

1.4.3 Dropping Unnecessary Columns

Since we will be focusing on domestic gross, we will drop worldwide gross data as well as any other unnecessary columns from our financial table.

```
[9]: budgets_df.drop(columns=['id', 'worldwide_gross'], inplace=True)
```

```
[10]: # Drop columns 'id' and 'worldwide_gross' which will not be used for analysis.

# del budgets_df['id']
# del budgets_df['worldwide_gross']
budgets_df.head()
```

```
[10]:
```

	release_date	movie	\
0	2009	Avatar	
1	2011	Pirates of the Caribbean: On Stranger Tides	
2	2019	Dark Phoenix	
3	2015	Avengers: Age of Ultron	
4	2017	Star Wars Ep. VIII: The Last Jedi	

	production_budget	domestic_gross
0	425000000	760507625
1	410600000	241063875
2	350000000	42762350
3	330600000	459005868
4	317000000	620181382

1.4.4 Dropping Missing Data Entries

We then need to address any rows where we have missing data for domestic gross.

```
[11]: # Replace all 0 values for domestic gross with NaN.

budgets_df['domestic_gross'] = budgets_df['domestic_gross'] \
    .map(lambda x: np.nan if x==0 else x)
```

```
[12]: # Return only rows where 'domestic_gross' is NOT NaN.

budgets_df = budgets_df[budgets_df['domestic_gross'].notna()]
```

1.4.5 Merging Basic Movie Data with Financial Data

This is where we will combine our IMDB dataset, where we have our genre and average rating data, with our The Numbers dataset which includes all of our financial data. In order to ensure that we are not incorrectly merging our financial data on different movies that have the same title, we will use the release year in conjunction with the title in our merge. We also want to make sure to keep only titles that have financial data, hence we will merge left onto our financial table.

```
[13]: # Join financial data from tn_movie_budgets with title and rating date
# from imdb_df.

merged_df = budgets_df.merge(imdb_df, how='left', left_on=['movie', \
```

```

right_on=['original_title', 'start_year'], \
merged_df.head()

```

```

[13]:
release_date      movie \
0      2009      Avatar
1      2011  Pirates of the Caribbean: On Stranger Tides
2      2019      Dark Phoenix
3      2015      Avengers: Age of Ultron
4      2017      Star Wars Ep. VIII: The Last Jedi

production_budget  domestic_gross    tconst \
0      425000000    760507625.0      NaN
1      410600000    241063875.0  tt1298650
2      350000000    42762350.0  tt6565702
3      330600000    459005868.0  tt2395427
4      317000000    620181382.0      NaN

primary_title \
0      NaN
1  Pirates of the Caribbean: On Stranger Tides
2      Dark Phoenix
3      Avengers: Age of Ultron
4      NaN

original_title  start_year  runtime_minutes \
0      NaN      NaN      NaN
1  Pirates of the Caribbean: On Stranger Tides    2011.0    136.0
2      Dark Phoenix    2019.0    113.0
3      Avengers: Age of Ultron    2015.0    141.0
4      NaN      NaN      NaN

genres  averagerating  numvotes
0      NaN      NaN      NaN
1  Action,Adventure,Fantasy    6.6  447624.0
2  Action,Adventure,Sci-Fi    6.0   24451.0
3  Action,Adventure,Sci-Fi    7.3  665594.0
4      NaN      NaN      NaN

```

```

[14]: merged_df.info()

```

```

<class 'pandas.core.frame.DataFrame'>
Int64Index: 5258 entries, 0 to 5257
Data columns (total 12 columns):
#   Column              Non-Null Count  Dtype
---

```



```

0  release_date      5258 non-null  int64
1  movie            5258 non-null  object
2  production_budget 5258 non-null  int64
3  domestic_gross    5258 non-null  float64
4  tconst           1330 non-null  object
5  primary_title     1330 non-null  object
6  original_title    1330 non-null  object
7  start_year        1330 non-null  float64
8  runtime_minutes   1325 non-null  float64
9  genres            1330 non-null  object
10 averagerating     1330 non-null  float64
11 numvotes          1330 non-null  float64
dtypes: float64(5), int64(2), object(5)
memory usage: 534.0+ KB

```

1.4.6 Checking for Missing Rating Data

Since it is possible that there was no matching data for each of the entries provided in the financial table, we need to remove any rows that do not have basic movie information data being matched from the IMDB dataset.

```

[15]: # Filter for rows that are not missing rating data.

merged_df = merged_df[merged_df['averagerating'].notna()] \
            .sort_values('release_date')

merged_df.head()

```

```

[15]:      release_date      movie  production_budget \
2530      2010      You Again      20000000
2501      2010  Vampires Suck      20000000
2467      2010  Why Did I Get Married Too?      20000000
2466      2010    The Last Song      20000000
2449      2010    Jackass 3D      20000000

      domestic_gross      tconst      primary_title \
2530      25702053.0  tt1414382      You Again
2501      36661504.0  tt1666186    Vampires Suck
2467      60095852.0  tt1391137  Why Did I Get Married Too?
2466      62950384.0  tt1294226    The Last Song
2449      117229692.0  tt1116184    Jackass 3D

      original_title  start_year  runtime_minutes \
2530      You Again      2010.0      105.0
2501    Vampires Suck      2010.0      82.0
2467  Why Did I Get Married Too?      2010.0      121.0
2466    The Last Song      2010.0      107.0
2449    Jackass 3D      2010.0      95.0

```

	genres	averagerating	numvotes
2530	Comedy,Family,Romance	5.8	46690.0
2501	Comedy	3.4	43984.0
2467	Comedy,Drama,Romance	4.6	8653.0
2466	Drama,Music,Romance	6.0	74914.0
2449	Action,Comedy,Documentary	7.0	53289.0

1.4.7 Checking and Dropping Duplicates

It is necessary to check for duplicates rows where movie and release date are matching. In order to prevent financial data from being matched with incorrect movies with the same title, we will drop the duplicates which have a lower number of rating votes.

```
[16]: # Drop duplicates based on movie title and release year

merged_df[merged_df.duplicated(subset=['movie', 'release_date'],
                                keep=False)].head()
```

```
[16]:      release_date      movie  production_budget  domestic_gross  \
2654      2010      The Tempest      20000000      277943.0
2653      2010      The Tempest      20000000      277943.0
1263      2010  The Bounty Hunter      45000000     67061228.0
1262      2010  The Bounty Hunter      45000000     67061228.0
1017      2010      Burlesque      55000000     39440655.0
```

	tconst	primary_title	original_title	start_year	\
2654	tt1683003	The Tempest	The Tempest	2010.0	
2653	tt1274300	The Tempest	The Tempest	2010.0	
1263	tt1472211	The Bounty Hunter	The Bounty Hunter	2010.0	
1262	tt1038919	The Bounty Hunter	The Bounty Hunter	2010.0	
1017	tt1586713	Burlesque	Burlesque	2010.0	

	runtime_minutes	genres	averagerating	numvotes
2654	131.0	Drama	7.8	94.0
2653	110.0	Comedy,Drama,Fantasy	5.4	7073.0
1263	NaN	None	6.3	29.0
1262	110.0	Action,Comedy,Romance	5.6	112444.0
1017	NaN	Drama	7.0	45.0

```
[17]: # Sort values by number of votes in preparation of dropping duplicates with
# lower vote count

merged_df.sort_values('numvotes', ascending=False, inplace=True)

# Drop duplicated movie with lower vote count
```

```
merged_df.drop_duplicates(subset=['movie', 'release_date'], inplace=True)

merged_df.head()
```

```
[17]:
```

	release_date	movie	production_budget	domestic_gross	\
139	2010	Inception	160000000	292576195.0	
10	2012	The Dark Knight Rises	275000000	448139099.0	
133	2014	Interstellar	165000000	188017894.0	
369	2012	Django Unchained	100000000	162805434.0	
26	2012	The Avengers	225000000	623279547.0	

	tconst	primary_title	original_title	start_year	\
139	tt1375666	Inception	Inception	2010.0	
10	tt1345836	The Dark Knight Rises	The Dark Knight Rises	2012.0	
133	tt0816692	Interstellar	Interstellar	2014.0	
369	tt1853728	Django Unchained	Django Unchained	2012.0	
26	tt0848228	The Avengers	The Avengers	2012.0	

	runtime_minutes	genres	averagerating	numvotes
139	148.0	Action,Adventure,Sci-Fi	8.8	1841066.0
10	164.0	Action,Thriller	8.4	1387769.0
133	169.0	Adventure,Drama,Sci-Fi	8.6	1299334.0
369	165.0	Drama,Western	8.4	1211405.0
26	143.0	Action,Adventure,Sci-Fi	8.1	1183655.0

```
[18]: # Verifying that duplicates have been eliminated.
```

```
len(merged_df[merged_df.duplicated(subset=['movie', 'release_date'], \
                                   keep=False)])
```

```
[18]: 0
```

1.4.8 Restricting Data to Relevant Years

```
[19]: # Restrict to last 5 years to keep focus on recent movies, but
# exclude 2020 due to covid impact.
```

```
merged_df = merged_df[(merged_df['release_date'] >= 2015) & \
                      (merged_df['release_date'] <= 2019)] \
               .sort_values('release_date')
```

1.4.9 Feature Engineering

Because our one of our KPIs is percentage profit, we need to create a column that displays this calculation from the domestic gross and production budget columns. The specific formula we will use to calculate percentage profit is $\$((\text{Domestic Gross} - \text{Production Budget}) / \text{Production Budget}) * 100 \%$.

```
[20]: # Create column for % profit based on the formula:
# (domestic_gross - production_budget) / production_budget

merged_df['% profit'] = ((merged_df['domestic_gross'] \
                           - merged_df['production_budget']) \
                          / merged_df['production_budget']) * 100
```

1.5 Data Analysis

```
[21]: # Set theme and style for plots.
sns.set_theme('talk')
sns.set_style('darkgrid')
```

1.5.1 FuncFormatter

Before we plot our visualizations, we will define a function to transform our dollar amounts into easier-to-read dollar amounts in millions.

```
[22]: # Define function for displaying large dollar amounts in millions.

def millions(x, pos):
    """Source: https://stackoverflow.com/questions/61330427/
    ↪set-y-axis-in-millions"""
    'The two args are the value and tick position'
    return '%1.0fM' % (x * 1e-6)

formatter = FuncFormatter(millions)
```

We have thoroughly prepared our data for visualization, and we can now return to our three questions for analysis:

1. What genres of movie are likely to succeed?
2. What genres cost more to produce?
3. How does production budget affect the success of a movie?

We will now proceed to plot our data to help us get a better sense of how each of these criteria translate to a movie's success.

1.5.2 Genre vs. Movie Success

To examine which genres have the highest KPIs, we will use bar plots. This enables us to clearly see the aggregate median values for each of our genres.

General Genre Data

```
[23]: # Split and explode entries to show one genre per row with repeated titles
# where necessary.

merged_df['genre_list'] = merged_df['genres'].str.split(',')
```

```

exploded_df = merged_df.explode('genre_list')

exploded_df.head()

```

```

[23]:
      release_date      movie  production_budget \
1657      2015      Concussion      35000000
1657      2015      Concussion      35000000
1657      2015      Concussion      35000000
3470      2015  The Second Best Exotic Marigold Hotel      10000000
3470      2015  The Second Best Exotic Marigold Hotel      10000000

      domestic_gross      tconst      primary_title \
1657      34531832.0  tt3322364      Concussion
1657      34531832.0  tt3322364      Concussion
1657      34531832.0  tt3322364      Concussion
3470      33078266.0  tt2555736  The Second Best Exotic Marigold Hotel
3470      33078266.0  tt2555736  The Second Best Exotic Marigold Hotel

      original_title  start_year  runtime_minutes \
1657      Concussion      2015.0      123.0
1657      Concussion      2015.0      123.0
1657      Concussion      2015.0      123.0
3470  The Second Best Exotic Marigold Hotel      2015.0      122.0
3470  The Second Best Exotic Marigold Hotel      2015.0      122.0

      genres  averagerating  numvotes      % profit  genre_list
1657  Biography,Drama,Sport      7.1      77576.0      -1.337623  Biography
1657  Biography,Drama,Sport      7.1      77576.0      -1.337623      Drama
1657  Biography,Drama,Sport      7.1      77576.0      -1.337623      Sport
3470      Comedy,Drama      6.6      28931.0      230.782660      Comedy
3470      Comedy,Drama      6.6      28931.0      230.782660      Drama

```

Genre vs. Rating

```

[24]: # Group by genre and calculate aggregate median sorted by rating
      # and return top 10 genres.

genre_rating_df = exploded_df.groupby('genre_list') \
      .median()[['averagerating','% profit']] \
      .sort_values('averagerating', ascending=False) \
      .head(10)
genre_rating_df.reset_index(inplace=True)

genre_rating_df

```

```

[24]:
      genre_list  averagerating      % profit
0      Musical      7.40      161.278340

```

1	History	7.20	2.056160
2	Sport	7.10	-1.337623
3	Biography	7.10	-5.573207
4	Western	6.90	3.814061
5	Drama	6.80	10.429731
6	Family	6.65	13.786461
7	Animation	6.60	23.805229
8	Music	6.55	38.294300
9	Documentary	6.50	-4.756800

Genre vs. Percent Profit

```
[25]: # Group by genre and calculate aggregate median sorted by % profit.
```

```
genre_profit_df = exploded_df.groupby('genre_list') \
    .median()[['averagerating', '% profit']] \
    .sort_values('% profit', ascending=False) \
    .head(10)

genre_profit_df.reset_index(inplace=True)
genre_profit_df
```

```
[25]:  genre_list  averagerating    % profit
0    Horror      5.80  178.555945
1   Musical      7.40  161.278340
2   Mystery      6.15  153.931679
3   Romance      6.40   42.099253
4    Comedy      6.30   38.617048
5     Music      6.55   38.294300
6 Animation      6.60   23.805229
7   Family      6.65   13.786461
8  Thriller      6.10   11.716700
9    Drama      6.80   10.429731
```

Plotting Genre vs. Success

```
[26]: # Plot genre vs averagerating on bar plot

fig, axes = plt.subplots(ncols=2, figsize=(16,6))

sns.barplot(data=genre_rating_df,
             x="genre_list", y="averagerating",
             ax=axes[0], palette='Blues_r', alpha=0.8)

axes[0].set_ylim([6, 7.5])

axes[0].set_title('Movie Genre vs. Rating')
axes[0].set_xlabel('Genre')
```

```


axes[0].set_ylabel('Average Rating')
axes[0].set_xticklabels(axes[0].get_xticklabels(), rotation=45, ha='right')

# Plot genre vs average percentage profit on a bar plot

sns.barplot(data=genre_profit_df,
            x="genre_list", y="% profit",
            ax=axes[1], palette='Greens_r', alpha=.8)

axes[1].set_title('Movie Genre vs. Percentage Profit')
axes[1].set_xlabel('Genre')
axes[1].set_ylabel('Average Percent Profit (%)')
axes[1].set_xticklabels(axes[1].get_xticklabels(), rotation=45, ha='right');

```



What genres of movie are likely to succeed? Our bar plot indicates that Musicals have a tendency to receive higher ratings with non-fictional genres including History, Biography and Sports following closely behind.

However, the most profitable genres by far appear to be Horror, Musicals, and Mystery.

The Musical genre appears to be a top performer in both cases, but otherwise, genre choice will depend on whether Microsoft's goal is to build a reputation for building good movies, or if it is purely to maximize profits in the most efficient use of its budget.

1.5.3 Genre vs. Production Costs

To examine which genres have the highest mean production costs, we will use bar plots. This plot will indicate clearly which genres on average cost the most to produce.

Genre vs Average Production Cost

```
[27]: # Group by genre and calculate aggregate median sorted by production budget.
      # Then create a high budget and low budget table.
```

```
high_budget_df = exploded_df.groupby('genre_list') \
                        .mean()[['production_budget']] \
                        .sort_values('production_budget', \
                                    ascending=False) \
                        .head(11)
high_budget_df.reset_index(inplace=True)
high_budget_df = high_budget_df[high_budget_df['genre_list'] != 'None']

low_budget_df = exploded_df.groupby('genre_list') \
                        .mean()[['production_budget']] \
                        .sort_values('production_budget', \
                                    ascending=False) \
                        .tail(10)
low_budget_df.reset_index(inplace=True)
low_budget_df = low_budget_df[low_budget_df['genre_list'] != 'None']

display(high_budget_df, low_budget_df)
```

	genre_list	production_budget
0	Musical	1.220000e+08
1	Adventure	1.102115e+08
2	Fantasy	1.086174e+08
3	Sci-Fi	1.059534e+08
4	Action	9.139595e+07
5	Western	9.000000e+07
6	Animation	8.890217e+07
7	Family	8.851333e+07
8	Comedy	4.551694e+07
9	History	3.823600e+07
10	Thriller	3.542267e+07

	genre_list	production_budget
0	Crime	3.451084e+07
1	Documentary	3.321000e+07
2	Drama	2.907010e+07
3	Biography	2.760667e+07
4	Sport	2.593333e+07
5	Romance	2.232653e+07
6	Music	1.945000e+07
7	Mystery	1.868558e+07
8	Horror	1.697647e+07
9	War	1.340000e+07

Plotting Genre vs. Production Budget


```
[28]: #Plot highest costing movies

fig, axes = plt.subplots(ncols=2, figsize=(16,6))

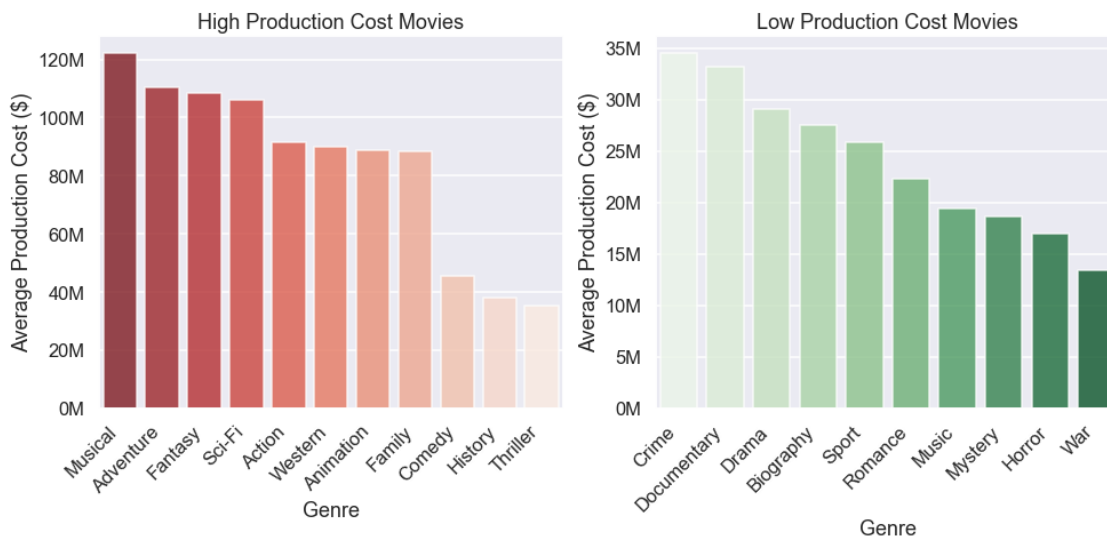
sns.barplot(data=high_budget_df,
            x="genre_list", y="production_budget",
            ax=axes[0], palette='Reds_r', alpha=.8)

axes[0].set_title('High Production Cost Movies')
axes[0].set_xlabel('Genre')
axes[0].set_ylabel('Average Production Cost ($)')
axes[0].set_xticklabels(axes[0].get_xticklabels(), rotation=45, ha='right')
axes[0].yaxis.set_major_formatter(formatter)

# Plot lowest costing movies

sns.barplot(data=low_budget_df,
            x="genre_list", y="production_budget",
            ax=axes[1], palette='Greens', alpha=.8)

axes[1].set_title('Low Production Cost Movies')
axes[1].set_xlabel('Genre')
axes[1].set_ylabel('Average Production Cost ($)')
axes[1].set_xticklabels(axes[1].get_xticklabels(), rotation=45, ha='right')
axes[1].yaxis.set_major_formatter(formatter);
```



What genres cost the most and least to produce? Musicals are by far the most costly genre of movie to produce, followed by Fantasy, Sci-Fi and Adventure. The lowest costing genres are War, Horror and Mystery.

1.5.4 Production Budget vs. Percentage Profit

Now that we have an idea of the impact of genre choice, we can begin to look at how production budget affects movie success.

Production Budget Data

```
[29]: # Create a copy of the dataframe to be used in our analysis of  
# budget vs. rating and budget vs. percent profit.
```

```
financial_df = merged_df  
financial_df.head()
```

```
[29]:      release_date      movie  production_budget \  
1657      2015      Concussion      35000000  
3470      2015  The Second Best Exotic Marigold Hotel      10000000  
1710      2015      Unfinished Business      35000000  
1163      2015      Run All Night      50000000  
1540      2015      Paul Blart: Mall Cop 2      38000000
```

```
      domestic_gross      tconst      primary_title \  
1657      34531832.0  tt3322364      Concussion  
3470      33078266.0  tt2555736  The Second Best Exotic Marigold Hotel  
1710      10219501.0  tt2358925      Unfinished Business  
1163      26461644.0  tt2199571      Run All Night  
1540      71091594.0  tt3450650      Paul Blart: Mall Cop 2
```

```
      original_title  start_year  runtime_minutes \  
1657      Concussion      2015.0      123.0  
3470  The Second Best Exotic Marigold Hotel      2015.0      122.0  
1710      Unfinished Business      2015.0      91.0  
1163      Run All Night      2015.0      114.0  
1540      Paul Blart: Mall Cop 2      2015.0      94.0
```

```
      genres  averagerating  numvotes  % profit \  
1657  Biography,Drama,Sport      7.1  77576.0  -1.337623  
3470      Comedy,Drama      6.6  28931.0  230.782660  
1710      Comedy,Drama      5.4  29004.0  -70.801426  
1163  Action,Drama,Thriller      6.6  94131.0  -47.076712  
1540      Action,Comedy,Crime      4.4  30828.0  87.083142
```

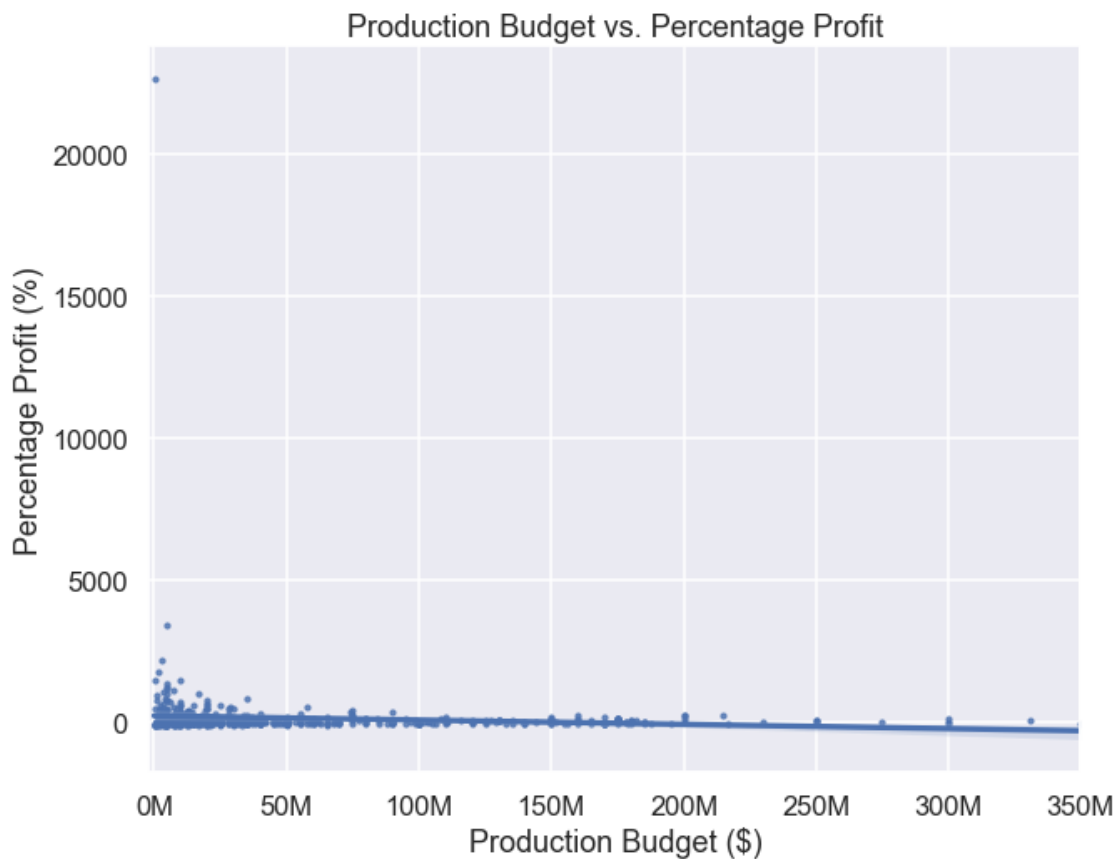
```
      genre_list  
1657  [Biography, Drama, Sport]  
3470      [Comedy, Drama]  
1710      [Comedy, Drama]  
1163  [Action, Drama, Thriller]  
1540      [Action, Comedy, Crime]
```

Plotting Production Budget vs. Success

```
[30]: # Plot budget vs % profit on regression scatter plot

fig, ax = plt.subplots(figsize=(10,8))
sns.regplot(x="production_budget", y="% profit",
            data=financial_df, ax=ax, fit_reg=True, scatter_kws={'s':10})

ax.set_xlim([-2000000, 350000000])
ax.set_title('Production Budget vs. Percentage Profit')
ax.set_xlabel('Production Budget ($)')
ax.set_ylabel('Percentage Profit (%)')
ax.xaxis.set_major_formatter(formatter);
```



1.5.5 Production Budget vs. Movie Success

We can see above that due to extreme outliers, our regression plot is distorted, showing a large cluster of points below 2,000% profit. Hence, we will restrict the y-axis to more clearly show our data points and regression line.

In this situation, it is appropriate to use a regression plot to display a marker for each of our movies, showing the relationship between each movie's production budget and its respective KPI.

```
[31]: # Plot budget vs averagerating on regression scatter plot

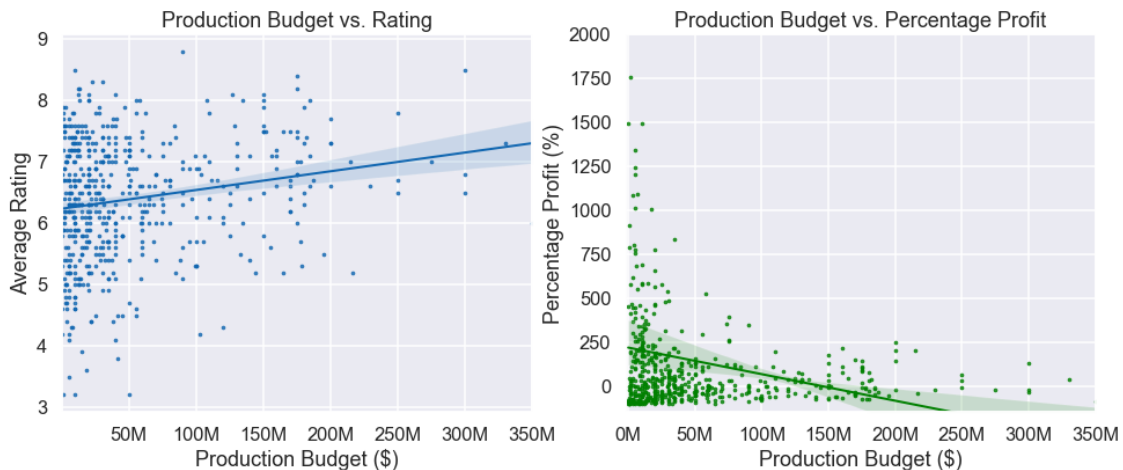
fig, axes = plt.subplots(ncols=2, figsize=(16,6))
sns.regplot(x="production_budget", y="averagerating", color='#1167b1',
            line_kws={"lw":2}, scatter_kws={'s':7}, data=financial_df,
            fit_reg=True, ax=axes[0])

axes[0].set_title('Production Budget vs. Rating')
axes[0].set_xlabel('Production Budget ($)')
axes[0].set_ylabel('Average Rating')
axes[0].xaxis.set_major_formatter(formatter)

# Plot and zoom in to budget vs percent profit to examine regression
# without severe outliers

sns.regplot(x="production_budget", y="% profit", color='green',
            line_kws={"lw":2}, data=financial_df, ax=axes[1],
            fit_reg=True, scatter_kws={'s':7})
axes[1].set_ylim([-150, 2000])
axes[1].set_xlim([-2000000, 350000000])

axes[1].set_title('Production Budget vs. Percentage Profit')
axes[1].set_xlabel('Production Budget ($)')
axes[1].set_ylabel('Percentage Profit (%)')
axes[1].xaxis.set_major_formatter(formatter);
```



How does production budget affect the success of a movie? We can see in our regression plot that as production budget increases, there is a positive trend in average rating. Conversely, we can see that profit percentages tend to decrease with larger budget movies.

In this case, it is important to note that the margin of error is skewed on the extremes of the budget. There are a much higher number of samples that are low budget, and fewer samples that are high budget.

Therefore, we can tell that a **higher budget is more likely to receive a higher rating, while likely to return a lower profit percentage**, but movie budget is not a definitive indicator of how successful a movie will be.

This plot gives us enough information now to make a conclusion on what Microsoft should keep in mind when creating its first films.

1.6 Conclusions

There is no clear-cut formula to creating a successful movie, and there will always be exceptions, even though a certain movie might not seem to fit criteria that have had a track record of success.

Horror and Mystery movies only require a low production budget, but have a track record of receiving high profit percentages. If profit percentage is a priority for Microsoft, either of these genres would be a good pick.

However, with the above analysis, we can see that although Musicals cost the most to produce, they do have a tendency to perform well as opposed to other genres. More specifically, Musicals were likely to receive better ratings, as well as have a higher profit percentage.

Under the assumption that production budget is not a major concern for Microsoft, we can conclude that it would be a safe choice to create movies under the Musical genre, despite the slightly negative correlation between production budget and profit percentage. By focusing on receiving higher ratings, Microsoft would be able to earn a strong reputation within the film industry, allowing it to more effectively market its future productions.

Some questions to consider for further analysis include the following:

1. What would be the most efficient allocation of production budget between cast, directors and writers and does how does this apply to the most successful genres? This would include an analysis of how much impact each of these roles tend to have on a genre's success.
2. Does the provided analysis apply when considering worldwide gross as opposed to just domestic gross? Although it would be best to start off focusing on a smaller audience, it could be in Microsoft's best interest to eventually increase production to a worldwide scale.

1.7 Appendix

1.7.1 Runtime vs. Movie Success

Another potentially useful metric to keep in mind is the movie runtime. This was not included as a main part of the analysis, since it is not as closely related to the other questions of analysis which dealt more with the impact of genre selection.

However, there does seem to be some correlation between runtime and production budget. As we increase both budget and runtime, they both tend to result in a higher rating but a lower profit percentage. It is also intuitive that longer movies cost more to produce due to the amount of additional editing time and wages that would need to be paid.

Runtime Data

```
[32]: # Eliminate entries missing data in 'runtime_minutes' to create
# runtime dataframe.
```

```
runtime_df = merged_df[merged_df['runtime_minutes'].notna()]
runtime_df
```

```
[32]:      release_date      movie \
1657      2015      Concussion
3470      2015  The Second Best Exotic Marigold Hotel
1710      2015      Unfinished Business
1163      2015      Run All Night
1540      2015  Paul Blart: Mall Cop 2
...      ...      ...
256      2019  How to Train Your Dragon: The Hidden World
125      2019      Alita: Battle Angel
3009      2019      Miss Bala
395      2019      Wonder Park
1462      2019      Long Shot
```

```
      production_budget  domestic_gross  tconst \
1657      35000000      34531832.0  tt3322364
3470      10000000      33078266.0  tt2555736
1710      35000000      10219501.0  tt2358925
1163      50000000      26461644.0  tt2199571
1540      38000000      71091594.0  tt3450650
...      ...      ...      ...
256      129000000      160791800.0  tt2386490
125      170000000      85710210.0  tt0437086
3009      15000000      14998027.0  tt5941692
395      100000000      45216793.0  tt6428676
1462      40000000      30202860.0  tt2139881
```

```
      primary_title \
1657      Concussion
3470  The Second Best Exotic Marigold Hotel
1710      Unfinished Business
1163      Run All Night
1540  Paul Blart: Mall Cop 2
...      ...
256  How to Train Your Dragon: The Hidden World
125      Alita: Battle Angel
3009      Miss Bala
395      Wonder Park
1462      Long Shot
```

```
      original_title  start_year  runtime_minutes \
1657      Concussion      2015.0      123.0
```

3470	The Second Best Exotic Marigold Hotel	2015.0	122.0
1710	Unfinished Business	2015.0	91.0
1163	Run All Night	2015.0	114.0
1540	Paul Blart: Mall Cop 2	2015.0	94.0
...
256	How to Train Your Dragon: The Hidden World	2019.0	104.0
125	Alita: Battle Angel	2019.0	122.0
3009	Miss Bala	2019.0	104.0
395	Wonder Park	2019.0	85.0
1462	Long Shot	2019.0	125.0

	genres	averagerating	numvotes	% profit	\
1657	Biography,Drama,Sport	7.1	77576.0	-1.337623	
3470	Comedy,Drama	6.6	28931.0	230.782660	
1710	Comedy,Drama	5.4	29004.0	-70.801426	
1163	Action,Drama,Thriller	6.6	94131.0	-47.076712	
1540	Action,Comedy,Crime	4.4	30828.0	87.083142	
...	
256	Action,Adventure,Animation	7.6	60769.0	24.644806	
125	Action,Adventure,Sci-Fi	7.5	88207.0	-49.582229	
3009	Action,Crime,Drama	5.5	3738.0	-0.013153	
395	Adventure,Animation,Comedy	5.7	3091.0	-54.783207	
1462	Comedy,Romance	7.2	12814.0	-24.492850	

	genre_list
1657	[Biography, Drama, Sport]
3470	[Comedy, Drama]
1710	[Comedy, Drama]
1163	[Action, Drama, Thriller]
1540	[Action, Comedy, Crime]
...	...
256	[Action, Adventure, Animation]
125	[Action, Adventure, Sci-Fi]
3009	[Action, Crime, Drama]
395	[Adventure, Animation, Comedy]
1462	[Comedy, Romance]

[546 rows x 14 columns]

```
[33]: # Group by runtime_minutes in order and calculate aggregate median for
# all columns.
```

```
runtime_df = runtime_df.sort_values('runtime_minutes')

runtime_df.head()
```

```
[33]:
```

	release_date	movie	production_budget	domestic_gross	\
193	2017	The Great Wall	150000000	45157105.0	
5157	2015	The Overnight	200000	1109808.0	
3613	2016	Kicks	10000000	150191.0	
4090	2016	Lights Out	5000000	67268835.0	
5196	2015	The Gallows	100000	22764410.0	

	tconst	primary_title	original_title	start_year	runtime_minutes	\
193	tt7535780	The Great Wall	The Great Wall	2017.0	72.0	
5157	tt3844362	The Overnight	The Overnight	2015.0	79.0	
3613	tt4254584	Kicks	Kicks	2016.0	80.0	
4090	tt4786282	Lights Out	Lights Out	2016.0	81.0	
5196	tt2309260	The Gallows	The Gallows	2015.0	81.0	

	genres	averagerating	numvotes	% profit	\
193	Documentary	6.5	24.0	-69.895263	
5157	Comedy,Mystery	6.1	14828.0	454.904000	
3613	Adventure,Drama	6.3	3789.0	-98.498090	
4090	Drama,Horror,Mystery	6.3	100650.0	1245.376700	
5196	Horror,Mystery,Thriller	4.2	17763.0	22664.410000	

	genre_list
193	[Documentary]
5157	[Comedy, Mystery]
3613	[Adventure, Drama]
4090	[Drama, Horror, Mystery]
5196	[Horror, Mystery, Thriller]

Runtime vs Movie Success To identify the relationship between runtime and movie success, we have used a regression plot to display a marker for each of the movies with an appropriate regression line for each KPI.

```
[34]: # Plot runtime minutes vs averagerating on regression scatter plot

fig, axes = plt.subplots(ncols=2, figsize=(16,6))

sns.regplot(x="runtime_minutes", y="averagerating", color='#1167b1',
            data=runtime_df, ax=axes[0], fit_reg=True, scatter_kws={'s':5},
            line_kws={"lw":2})

axes[0].set_title('Runtime vs. Rating')
axes[0].set_xlabel('Runtime (minutes)')
axes[0].set_ylabel('Average Rating')

# Plot runtime minutes vs percent profit on regression scatter plot

sns.regplot(x="runtime_minutes", y="% profit", color='green',
```



```

data=runtime_df, ax=axes[1], fit_reg=True, scatter_kws={'s':5},
line_kws={"lw":1})

axes[1].set_title('Runtime vs. Percentage Profit')
axes[1].set_xlabel('Runtime (minutes)')
axes[1].set_ylabel('Percentage Profit (%)')
axes[1].set_ylim([-150, 2000])

plt.tight_layout();

```

